

# Graded Practical Session

December 19, 2023

## 1 Exercise #1: Supervised Learning

In a regression problem, we have to predict a continuous dependent variable, like a price, from independent variables. We will use the dataset Auto MPG [site](#) and build some models to predict the energy efficiency (MPG) of vehicles of the end of 1970's and the beginning of 1980's. The independent variables (attributes) in this dataset are : # of cylinders, displacement, horsepower, weight, acceleration, model year, origin, car name. 1. Download the dataset at: [dataset](#) 2. Read the data with pandas. 3. Clean the data by possibly removing rows with unknown values (dropna with pandas). 4. Visualize the data with the most adapted techniques to have a glance about the correlation of some pairs of attributes. 5. Divide the data into a training set and a test set (80% training, 20% test). 6. Normalize the data (preprocessing by normalization). 7. Train a linear regression model first, then train a Deep Neuron Network with two dense layers with relu activation functions. For the neuron network, use the optimizer Adam, and test the loss functions mean\_absolute\_error and mean\_squared\_error. 8. Compare the results of the different models (linear regression and MLP) on the test set.

For steps 2 to 8, write Python codes.

```
[1]: # 2. Read the data with pandas.

import pandas as pd

# Define the column names
column_names = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
                'acceleration', 'model_year', 'origin', 'car_name']

# Read the data
data = pd.read_csv('auto-mpg.data', delim_whitespace=True, names=column_names)

data.head()
```

```
[1]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	\
0	18.0	8	307.0	130.0	3504.0	12.0	70	
1	15.0	8	350.0	165.0	3693.0	11.5	70	
2	18.0	8	318.0	150.0	3436.0	11.0	70	
3	16.0	8	304.0	150.0	3433.0	12.0	70	
4	17.0	8	302.0	140.0	3449.0	10.5	70	

	origin	car_name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

```
[2]: # 3. Clean the data by possibly removing rows with unknown values (dropna with
      ↪ pandas).

      # Replace '?' by NaN values
      ↪ (horsepower      Feature      Continuous      missing
      ↪ : yes)
data['horsepower'].replace('?', pd.NA, inplace=True)

      # Drop rows with NaN values
data.dropna(inplace=True)

data.head()
```

```
[2]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model_year  \
0   18.0         8         307.0        130.0   3504.0         12.0         70
1   15.0         8         350.0        165.0   3693.0         11.5         70
2   18.0         8         318.0        150.0   3436.0         11.0         70
3   16.0         8         304.0        150.0   3433.0         12.0         70
4   17.0         8         302.0        140.0   3449.0         10.5         70
```

	origin	car_name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

```
[3]: # 4. Visualize the data with the most adapted techniques to have a glance about
      ↪ the correlation of some pairs of attributes.

import seaborn as sns
import matplotlib.pyplot as plt


      # We drop the car_name column because it is not useful for the visualization
data_num = data.drop(['car_name'], axis=1)
```

```
[4]: # Pairplot to visualize the correlation between the attributes
sns.pairplot(data_num)
```


/Users/thibaultchaussou/miniconda3/envs/AI53/lib/python3.8/site-packages/seaborn/axisgrid.py:118: UserWarning: The figure layout has changed to

```
tight
self._figure.tight_layout(*args, **kwargs)
```

```
[4]: <seaborn.axisgrid.PairGrid at 0x15e034c40>
```



```
[5]: # Heatmap of the correlation between the attributes
plt.figure(figsize=(10, 10))
sns.heatmap(data_num.corr(), annot=True, cmap='coolwarm')
plt.show()
```



The most correlated attributes are : - mpg and model\_year - mpg and origin - cylinders and displacement - cylinders and horsepower - cylinders and weight - displacement and horsepower - displacement and weight - horsepower and weight

[6]: *# 5. Divide the data into a training set and a test set (80% training, 20% test).*

```
from sklearn.model_selection import train_test_split

# Split the data into X and y
X = data.drop(['mpg', 'car_name', 'origin'], axis=1)
y = data['mpg'] # Target
```

```
# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2) # We
→ can use also random_state=??
```

[7]: *# 6. Normalize the data (preprocessing by normalization).*

```
from sklearn.preprocessing import StandardScaler

# Normalize the data, with the StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```


[8]: *# 7. Train a linear regression model first, then train a Deep Neuron Network*  
*→ with two dense layers with relu activation functions.*

```
from sklearn.linear_model import LinearRegression

# Training the model linear regression
lr = LinearRegression()
lr.fit(X_train_scaled, y_train)

# Predictions
y_pred_lr = lr.predict(X_test_scaled)

# Plot the predictions
plt.figure(figsize=(10, 10))
plt.scatter(y_test, y_pred_lr)
plt.xlabel('Valeurs réelles')
plt.ylabel('Valeurs prédites')
plt.show()
```



```
[9]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    Dense(64, activation='relu'),
    Dense(1)
])

model.compile(optimizer='adam', loss='mean_squared_error')
model.fit(X_train_scaled, y_train, epochs=100)

y_pred_m = model.predict(X_test_scaled)
```

```
plt.figure(figsize=(10, 10))
plt.scatter(y_test, y_pred_m)
plt.xlabel('Valeurs réelles')
plt.ylabel('Valeurs prédites')
plt.show()
```

Epoch 1/100

```
2023-12-19 18:42:51.868674: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M1 Pro
2023-12-19 18:42:51.868728: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 16.00 GB
2023-12-19 18:42:51.868743: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 5.33 GB
2023-12-19 18:42:51.868818: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:303]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2023-12-19 18:42:51.868851: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:269]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)
```

1/10 [==>...] - ETA: 3s - loss: 552.6580

```
2023-12-19 18:42:52.244746: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:114]
Plugin optimizer for device_type GPU is enabled.
```

10/10 [=====] - 0s 14ms/step - loss: 588.2451

Epoch 2/100

10/10 [=====] - 0s 6ms/step - loss: 576.4590

Epoch 3/100

10/10 [=====] - 0s 6ms/step - loss: 564.2844

Epoch 4/100

10/10 [=====] - 0s 6ms/step - loss: 548.8628

Epoch 5/100

10/10 [=====] - 0s 6ms/step - loss: 532.4020

Epoch 6/100

10/10 [=====] - 0s 6ms/step - loss: 511.8532

Epoch 7/100

10/10 [=====] - 0s 6ms/step - loss: 491.9503

Epoch 8/100

10/10 [=====] - 0s 6ms/step - loss: 474.2872

Epoch 9/100

10/10 [=====] - 0s 6ms/step - loss: 469.3499

Epoch 10/100

10/10 [=====] - 0s 6ms/step - loss: 460.5447

Epoch 11/100  
10/10 [=====] - 0s 6ms/step - loss: 448.7555  
Epoch 12/100  
10/10 [=====] - 0s 6ms/step - loss: 434.0291  
Epoch 13/100  
10/10 [=====] - 0s 6ms/step - loss: 417.3637  
Epoch 14/100  
10/10 [=====] - 0s 6ms/step - loss: 400.4912  
Epoch 15/100  
10/10 [=====] - 0s 6ms/step - loss: 385.8094  
Epoch 16/100  
10/10 [=====] - 0s 6ms/step - loss: 371.6675  
Epoch 17/100  
10/10 [=====] - 0s 6ms/step - loss: 357.7856  
Epoch 18/100  
10/10 [=====] - 0s 6ms/step - loss: 339.0591  
Epoch 19/100  
10/10 [=====] - 0s 6ms/step - loss: 324.7419  
Epoch 20/100  
10/10 [=====] - 0s 6ms/step - loss: 308.6317  
Epoch 21/100  
10/10 [=====] - 0s 6ms/step - loss: 287.8429  
Epoch 22/100  
10/10 [=====] - 0s 6ms/step - loss: 275.0320  
Epoch 23/100  
10/10 [=====] - 0s 6ms/step - loss: 255.3445  
Epoch 24/100  
10/10 [=====] - 0s 6ms/step - loss: 242.9864  
Epoch 25/100  
10/10 [=====] - 0s 6ms/step - loss: 223.8944  
Epoch 26/100  
10/10 [=====] - 0s 6ms/step - loss: 209.1921  
Epoch 27/100  
10/10 [=====] - 0s 6ms/step - loss: 194.1037  
Epoch 28/100  
10/10 [=====] - 0s 6ms/step - loss: 183.2771  
Epoch 29/100  
10/10 [=====] - 0s 6ms/step - loss: 167.0127  
Epoch 30/100  
10/10 [=====] - 0s 6ms/step - loss: 151.4143  
Epoch 31/100  
10/10 [=====] - 0s 6ms/step - loss: 139.3593  
Epoch 32/100  
10/10 [=====] - 0s 6ms/step - loss: 129.1301  
Epoch 33/100  
10/10 [=====] - 0s 6ms/step - loss: 125.4392  
Epoch 34/100  
10/10 [=====] - 0s 6ms/step - loss: 113.9862



Epoch 35/100  
10/10 [=====] - 0s 6ms/step - loss: 99.2163  
Epoch 36/100  
10/10 [=====] - 0s 6ms/step - loss: 85.7128  
Epoch 37/100  
10/10 [=====] - 0s 6ms/step - loss: 79.2580  
Epoch 38/100  
10/10 [=====] - 0s 6ms/step - loss: 72.7921  
Epoch 39/100  
10/10 [=====] - 0s 6ms/step - loss: 64.3281  
Epoch 40/100  
10/10 [=====] - 0s 6ms/step - loss: 55.8739  
Epoch 41/100  
10/10 [=====] - 0s 6ms/step - loss: 48.6070  
Epoch 42/100  
10/10 [=====] - 0s 6ms/step - loss: 42.9457  
Epoch 43/100  
10/10 [=====] - 0s 6ms/step - loss: 40.9757  
Epoch 44/100  
10/10 [=====] - 0s 6ms/step - loss: 35.5096  
Epoch 45/100  
10/10 [=====] - 0s 6ms/step - loss: 32.1649  
Epoch 46/100  
10/10 [=====] - 0s 6ms/step - loss: 29.1743  
Epoch 47/100  
10/10 [=====] - 0s 6ms/step - loss: 27.2038  
Epoch 48/100  
10/10 [=====] - 0s 6ms/step - loss: 26.7095  
Epoch 49/100  
10/10 [=====] - 0s 6ms/step - loss: 23.7697  
Epoch 50/100  
10/10 [=====] - 0s 6ms/step - loss: 23.8458  
Epoch 51/100  
10/10 [=====] - 0s 6ms/step - loss: 22.5669  
Epoch 52/100  
10/10 [=====] - 0s 6ms/step - loss: 21.9943  
Epoch 53/100  
10/10 [=====] - 0s 6ms/step - loss: 22.6254  
Epoch 54/100  
10/10 [=====] - 0s 6ms/step - loss: 21.7633  
Epoch 55/100  
10/10 [=====] - 0s 6ms/step - loss: 18.5290  
Epoch 56/100  
10/10 [=====] - 0s 6ms/step - loss: 17.8456  
Epoch 57/100  
10/10 [=====] - 0s 6ms/step - loss: 16.0783  
Epoch 58/100  
10/10 [=====] - 0s 6ms/step - loss: 17.5194

Epoch 59/100  
10/10 [=====] - 0s 6ms/step - loss: 16.9029  
Epoch 60/100  
10/10 [=====] - 0s 6ms/step - loss: 15.4360  
Epoch 61/100  
10/10 [=====] - 0s 6ms/step - loss: 15.0719  
Epoch 62/100  
10/10 [=====] - 0s 6ms/step - loss: 15.6680  
Epoch 63/100  
10/10 [=====] - 0s 6ms/step - loss: 15.2935  
Epoch 64/100  
10/10 [=====] - 0s 6ms/step - loss: 18.5763  
Epoch 65/100  
10/10 [=====] - 0s 6ms/step - loss: 15.7739  
Epoch 66/100  
10/10 [=====] - 0s 6ms/step - loss: 18.5788  
Epoch 67/100  
10/10 [=====] - 0s 6ms/step - loss: 15.3381  
Epoch 68/100  
10/10 [=====] - 0s 6ms/step - loss: 15.1738  
Epoch 69/100  
10/10 [=====] - 0s 6ms/step - loss: 17.9011  
Epoch 70/100  
10/10 [=====] - 0s 6ms/step - loss: 15.1687  
Epoch 71/100  
10/10 [=====] - 0s 6ms/step - loss: 14.8854  
Epoch 72/100  
10/10 [=====] - 0s 6ms/step - loss: 13.9522  
Epoch 73/100  
10/10 [=====] - 0s 6ms/step - loss: 14.3290  
Epoch 74/100  
10/10 [=====] - 0s 6ms/step - loss: 15.1516  
Epoch 75/100  
10/10 [=====] - 0s 6ms/step - loss: 15.5304  
Epoch 76/100  
10/10 [=====] - 0s 6ms/step - loss: 13.5889  
Epoch 77/100  
10/10 [=====] - 0s 6ms/step - loss: 13.8792  
Epoch 78/100  
10/10 [=====] - 0s 6ms/step - loss: 14.3180  
Epoch 79/100  
10/10 [=====] - 0s 6ms/step - loss: 13.8890  
Epoch 80/100  
10/10 [=====] - 0s 6ms/step - loss: 14.2448  
Epoch 81/100  
10/10 [=====] - 0s 6ms/step - loss: 14.0660  
Epoch 82/100  
10/10 [=====] - 0s 6ms/step - loss: 14.6394

```


Epoch 83/100
10/10 [=====] - 0s 6ms/step - loss: 14.0566
Epoch 84/100
10/10 [=====] - 0s 6ms/step - loss: 13.7054
Epoch 85/100
10/10 [=====] - 0s 6ms/step - loss: 15.7645
Epoch 86/100
10/10 [=====] - 0s 6ms/step - loss: 15.9438
Epoch 87/100
10/10 [=====] - 0s 6ms/step - loss: 15.7292
Epoch 88/100
10/10 [=====] - 0s 6ms/step - loss: 17.0639
Epoch 89/100
10/10 [=====] - 0s 6ms/step - loss: 13.6736
Epoch 90/100
10/10 [=====] - 0s 6ms/step - loss: 13.0268
Epoch 91/100
10/10 [=====] - 0s 6ms/step - loss: 13.1386
Epoch 92/100
10/10 [=====] - 0s 6ms/step - loss: 13.4674
Epoch 93/100
10/10 [=====] - 0s 6ms/step - loss: 13.7303
Epoch 94/100
10/10 [=====] - 0s 6ms/step - loss: 14.0973
Epoch 95/100
10/10 [=====] - 0s 6ms/step - loss: 19.4503
Epoch 96/100
10/10 [=====] - 0s 6ms/step - loss: 20.3662
Epoch 97/100
10/10 [=====] - 0s 6ms/step - loss: 18.3522
Epoch 98/100
10/10 [=====] - 0s 6ms/step - loss: 19.4924
Epoch 99/100
10/10 [=====] - 0s 6ms/step - loss: 15.2722
Epoch 100/100
10/10 [=====] - 0s 6ms/step - loss: 15.2105
3/3 [=====] - 0s 8ms/step

```

```

2023-12-19 18:42:58.802134: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:114]
Plugin optimizer for device_type GPU is enabled.

```



[10]: # 8. Compare the results of the different models (linear regression and MLP) on the test set.

```
from sklearn.metrics import mean_squared_error, mean_absolute_error

print('Linear Regression')
print('MSE :', mean_squared_error(y_test, y_pred_lr))
print('MAE :', mean_absolute_error(y_test, y_pred_lr))

print('\nMLP')
print('MSE :', mean_squared_error(y_test, y_pred_m))
print('MAE :', mean_absolute_error(y_test, y_pred_m))
```

Linear Regression

MSE : 12.62476754954062  
MAE : 2.7425887154000312


MLP

MSE : 13.962086354786473  
MAE : 2.6886178125309037

We note that the linear regression model performs slightly less well than the neural network model. But the computation time of the neural network is much longer than that of linear regression. For this problem, linear regression is more appropriate.

## 2 Exercise #2: Unsupervised Learning

We assume the following cloud of 20 points randomly chosen in the interval  $[-10;10]$  :



Program a Python code that: 1. generates the random points, 2. by either using the template of the graded practice session about K-Means, program the K- Means algorithm with N centers, or use it from a Python library (scikit-learn for example), in order to find the best possible clustering of the previous data consisting of 20 random points ; we assume that the coordinates of the cluster centers in K-Means are also chosen in the interval  $[-10;10]$ . 3. Finally, program the algorithm that consists in executing multiple K-Means with different numbers of centers and displaying the best result.

```
[11]: # 1. Generate the random points

import numpy as np

# Number of points per group
n_points = [4, 7, 9]

# All the centers of the groups
centers = [(0, 0), (5, 5), (-5, -5)]

# Generate points around each center
```


```

points = []
k = 0
for center in centers:
    x_points = np.random.normal(center[0], 1, n_points[k]) # 1 est l'écart-type
    y_points = np.random.normal(center[1], 1, n_points[k])
    points.append(np.column_stack((x_points, y_points)))
    k += 1

# Merge all the points
points = np.vstack(points)

# Plot the points
plt.figure(figsize=(10, 10))
plt.scatter(points[:, 0], points[:, 1])
plt.title("Points aléatoires")
plt.xlabel("X")
plt.ylabel("Y")
plt.show()

```



```
[12]: # 2. Program the algorithm that consists in executing multiple K-Means with  
      ↪ different numbers of centers and displaying the best result.
```


```
from sklearn.cluster import KMeans
```

```
def range_kmeans(nb_cluster):  
    # Run K-Means with different numbers of centers  
    inertia = []  
    K = [i for i in range(1, nb_cluster + 1)]  
    for k in K:  
        kmeans_fct = KMeans(n_clusters=k, n_init=10)
```

```
kmeans_fct.fit(points)
inertia.append(kmeans_fct.inertia_)

# Plot the results
plt.figure(figsize=(10, 10))
plt.plot(K, inertia, 'x-')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.title('Inertia vs. Number of clusters')
plt.show()
```

```
range_kmeans(10)
```






We note that the optimal number of clusters is 3. Because the inertia decreases significantly from 1 to 3, then it decreases slightly from 3 to 10.

```
[13]: # 3. Show the best result



nb_best_cluster = 3
kmeans = KMeans(n_clusters=nb_best_cluster, n_init=10)
kmeans.fit(points)

# Plot the points
plt.figure(figsize=(10, 10))
plt.scatter(points[:, 0], points[:, 1], c=kmeans.labels_)
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], c='r',
↪marker='+', s=200)
plt.title("Clusters")
plt.xlabel("X")
plt.ylabel("Y")
plt.show()
```



### 3 Exercise #3: Reinforcement Learning

In the problem of the arrows,  $n$  arrows are positioned vertically and oriented upwards or downwards. We want to swap the orientation of each arrow (see the figures below).



We can modify a position with the two following cases: - by swapping a sequence of three adjacent arrows having the same orientation (all upwards or all downwards) - by swapping two adjacent arrows having opposite orientations (one is upwards and the other is downwards). Questions: 1. Explain how you can define the MDP for this problem. Describe as clearly as possible how do you represent a state. 2. Program the environment using the template available on Moodle. You can modify it as you wish. 3. Give the optimal policy obtained by Q-Iteration for the initial state above.

#### 1. MDP Definition

We need to define the following elements (States, Actions, Rewards, Transitions, Termination Criterion): **States:**

A state can be represented by an array with the following letters +1 when the arrow points up or -1 when the arrow points down, where each element represents an arrow oriented upwards ('↑') or downwards ('↓'). For example, the initial state of our problem can be represented by [1, -1, 1, 1, 1, -1, 1, -1, -1, 1].

#### Actions:

The two types of actions are: - Inverting a sequence of three consecutive arrows of the same orientation. - Inverting two consecutive arrows of opposite orientations.

#### Transitions:

State transitions are deterministic in this case, as one is either up or down (no in-between possible). Each action leads to a new state in a predictable manner, if down → up, if up → down.

#### Rewards:

The reward is used to encourage reaching the state where all arrows are oriented in the opposite direction. We can define the reward as follows: - -1 for each action - 0 for the final state or more

#### Termination Criterion:

The game ends when all arrows have been turned in the opposite direction.

[14]: # 2. Program the environment using the template available on Moodle. You can   
→ modify it as you wish.

```

import itertools

class ArrowProblem:
    def __init__(self, n_arrows):
        self.n_arrows = n_arrows
        self.states = self.generate_states()

    def generate_states(self):
        return list(itertools.product([-1, 1], repeat=self.n_arrows))

    def state_space(self):
        return self.states

    def action_space(self, s):
        actions = []
        for i in range(self.n_arrows - 1):
            if s[i] != s[i + 1]: # Reverse two adjacent arrows in opposite
→directions
                actions.append((i, 2))
        for i in range(self.n_arrows - 2):
            if s[i] == s[i + 1] == s[i + 2]: # Reverse three adjacent arrows
→with the same orientation
                actions.append((i, 3))
        return actions

    def T(self, s, a):
        new_state = list(s)
        if a[1] == 2: # Reverse two arrows
            new_state[a[0]] = -s[a[0]]
            new_state[a[0] + 1] = -s[a[0] + 1]
        elif a[1] == 3: # Reverse three arrows
            new_state[a[0]] = -s[a[0]]
            new_state[a[0] + 1] = -s[a[0] + 1]
            new_state[a[0] + 2] = -s[a[0] + 2]
        return tuple(new_state)

    def R(self, s, a):
        return -1 # Simple -1 reward for each action

```

[15]: # 3. Give the optimal policy obtained by Q-Iteration for the initial state above.

```

import numpy as np
import copy

```

```

class QIteration:
    def __init__(self, mdp, gamma=0.9, precision=1e-13):
        self.MDP = mdp
        self.gamma = gamma
        self.precision = precision
        self.Q = dict()

        nba = 0
        for s in self.MDP.state_space():
            nba = max(nba, len(self.MDP.action_space(s)))
        for s in self.MDP.state_space():
            self.Q[s] = np.zeros(nba)

    def run(self, N=-1):
        l = 0
        while l != N:
            l += 1
            oldQ = copy.deepcopy(self.Q)

            for s in self.MDP.state_space():
                for a_idx, a in enumerate(self.MDP.action_space(s)):
                    next_state = self.MDP.T(s, a)
                    self.Q[s][a_idx] = self.MDP.R(s, a) + self.gamma * np.
→max(self.Q[next_state])

            if N == -1 and all(np.linalg.norm(oldQ[s] - self.Q[s]) < self.
→precision for s in self.MDP.state_space()):
                print('QIteration has stopped after', l, 'iterations!')
                break

    def find_optimal_policy(self):
        policy = {}
        for s in self.MDP.state_space():
            possible_actions = self.MDP.action_space(s)
            if possible_actions: # Vérifier s'il y a des actions possibles
                q_values = [self.Q[s][i] for i, _ in enumerate(possible_actions)]
                a_idx = np.argmax(q_values)
                policy[s] = possible_actions[a_idx]
            else:
                policy[s] = None
        return policy

```

```

[16]: # Create an instance of the arrow problem
arrow_problem = ArrowProblem(n_arrows=3)

# Create an instance of Q-Iteration
q_iteration = QIteration(arrow_problem)

```

```

# Run the Q-Iteration
q_iteration.run()

# Find the optimal policy
optimal_policy = q_iteration.find_optimal_policy()

# Print the optimal policy for some states
for state in arrow_problem.state_space()[:5]:
    print(f"État: {state}, Meilleure action: {optimal_policy[state]}")

```

QIteration has stopped after 3 iterations!

État: (-1, -1, -1), Meilleure action: (0, 3)

État: (-1, -1, 1), Meilleure action: (1, 2)

État: (-1, 1, -1), Meilleure action: (0, 2)

État: (-1, 1, 1), Meilleure action: (0, 2)

État: (1, -1, -1), Meilleure action: (0, 2)

1. (-1, -1, -1), Best Action: (0, 3): State: All arrows pointing down (-1). Action: (0, 3) means to start inverting arrows from index 0 and affect 3 consecutive arrows. In this state, it would flip all arrows to point upwards.
2. (-1, -1, 1), Best Action: (1, 2): State: The first two arrows point down, and the last one points up. Action: (1, 2) means to start inverting arrows from index 1 and affect 2 consecutive arrows. This would flip the second and third arrow.
3. (-1, 1, -1), Best Action: (0, 2): State: The first arrow points down, the second up, and the third down. Action: (0, 2) indicates inverting the first two adjacent arrows that have opposite orientations.
4. (-1, 1, 1), Best Action: (0, 2): State: The first arrow points down, the other two point up. Action: (0, 2) again indicates inverting the first two adjacent arrows with opposite orientations.
5. (1, -1, -1), Best Action: (1, 2): State: The first arrow points up, the other two point down. Action: (1, 2) indicates inverting the second and third arrows.

[17]: *# For our initial problem with 11 arrows, we can display the optimal policy for*  
*→each state:*

```

arrow_problem = ArrowProblem(n_arrows=11)

# Run the Q-Iteration
q_iteration = QIteration(arrow_problem)

# Run the Q-Iteration
q_iteration.run()

# Find the optimal policy
optimal_policy = q_iteration.find_optimal_policy()

# Print the optimal policy for some states
for state in arrow_problem.state_space()[:5]:

```

```
print(f"État: {state}, Meilleure action: {optimal_policy[state]}")
```

QIteration has stopped after 3 iterations!

État: (-1, -1, -1, -1, -1, -1, -1, -1, -1, -1), Meilleure action: (0, 3)

État: (-1, -1, -1, -1, -1, -1, -1, -1, -1, 1), Meilleure action: (9, 2)

État: (-1, -1, -1, -1, -1, -1, -1, -1, 1, -1), Meilleure action: (8, 2)

État: (-1, -1, -1, -1, -1, -1, -1, -1, 1, 1), Meilleure action: (8, 2)

État: (-1, -1, -1, -1, -1, -1, -1, 1, -1, -1), Meilleure action: (7, 2)

```
[18]: # Find the optimal policy for the initial state represented by [1, -1, 1, 1, 1, ↵
      ↵-1, 1, -1, -1, -1, 1, 1].
```

```
print(f"État initial : [1, -1, 1, 1, 1, -1, 1, -1, -1, 1, 1]")
print(f"Meilleure action : {optimal_policy[(1, -1, 1, 1, 1, -1, 1, -1, -1, 1, ↵
      ↵1)]}")
```

État initial : [1, -1, 1, 1, 1, -1, 1, -1, -1, 1, 1]

Meilleure action : (0, 2)

We perform the action recommended by the Q-Iteration algorithm and look for the optimal policy for the following state:  $[1, -1, 1, 1, 1, -1, 1, -1, -1, 1, 1] \rightarrow [1, 1, -1, -1, -1, 1, -1, 1, 1, 1, -1, -1]$  ETC.